



Indian Institute of Technology, Kharagpur
Department of Computer Science and Engineering

CS39002 : OPERATING SYSTEMS LAB

ASSIGNMENT 3: HANDS-ON EXPERIENCE OF USING SHARED MEMORY

27th Feb, 2023

Group Number : 24

Shivam Raj
20CS10056

Jatin Gupta
20CS10087

Kushaz Sehgal
20CS30030

Rushil Venkateswar
20CS30045

Contents

1. Data Structures	2
1.1 Description of the DNode class	2
1.2 Description of the Graph class	2
2. The Main Process	3
2.1 Description	3
2.2 Implementation	3
3. The Producer Process	3
3.1 Description	3
3.2 Implementation	3
4. The Consumer Process	3
4.1 Description	3
4.2 Implementation	3
5. Optimization on Dijkstra	4
5.1 Description	4
5.2 Implementation	4

1. Data Structures

1.1 Description of the DNode class

Data members :

- `int value` : the ID of the node
- `size_t next` : the index of the next node in the pool

1.2 Description of the Graph class

Data members :

- `DNode pool[MAX_LIST]` : central pool of nodes of size `MAX_LIST`. Each edge in the graph uses 2 nodes out of the pool
- `size_t node_to_head[MAX_NODE]` : stores the index of the head of the adjacency list of each node
- `size_t node_to_tail[MAX_NODE]` : stores the index of the tail of the adjacency list of each node
- `int degree[MAX_NODE]` : stores the degree of each node
- `size_t npool` : pointer to the next free index in the pool
- `int num_of_nodes` : total number of nodes in our Graph

Member functions :

- `void init()` : initializes the data members of Graph class
- `DNode * dnode_alloc()` : returns a pointer to a node if there are unused nodes in the pool
- `DNode * dnode(size_t index)` : returns a pointer to a node in the pool based on the index
- `DNode * dnode_next(const DNode *node)` : returns a pointer to the next node to the current node
- `void add_dnode(DNode *node, int a, int b)` : adds node b to the pool and sets b to be the tail of the adjacency list of node a
- `void dnode_push(int a, int b)` : allocate the nodes a and b in the pool and adds a to b's adjacency list, and vice versa
- `void dijkstra(vector<int> sources, vector<int> &dist, vector<int> &parent)` : multi-source dijkstra's algorithm called by each consumer process
- `void optimized_dijkstra(vector<int> sources, vector<int> &dist, vector<int> &parent)` : multi-source dijkstra's algorithm implemented level-wise
- `void propagate_new_nodes(vector<int> &dist, vector<int> &parent, vector<int> new_nodes)` : updates the non-source nodes (for a particular consumer) among the new nodes added

2. The Main Process

2.1 Description

Loads a graph into a shared memory. Note that the graph is dynamic - new nodes and new edges will be added later. It then spawns the producer and consumer processes. For simplicity, assume all edges of the graph to have weight 1.

2.2 Implementation

The first step is the creation of a Graph class which would store the dataset in the shared memory segment. This Graph class has a central pool of nodes (stored as an array). For each edge in the dataset, we create two nodes in the graph (since the graph is undirected) and update the `node_to_head` and `node_to_tail` arrays.

The main process reads the data and creates the graph in the shared memory. Once that is done, it forks and the child calls the producer process. Similarly, we fork 10 more times and call the consumer process in each child process.

3. The Producer Process

3.1 Description

Wakes up every 50 seconds, and updates the graph in the following way: A number m in the range $[10, 30]$ is chosen randomly, and m new nodes are added to the graph. For each newly added node, select a number k in the range $[1, 20]$ at random; the new node will connect to k existing nodes. The probability of a new node connecting to an existing node of degree d is proportional to d .

3.2 Implementation

The producer process launches an infinite loop in which it sleeps for 50 seconds and then selects a random value m . For each of these m new nodes to be added to the graph, we use the `discrete_distribution` function of `random` class to ascertain which nodes are to be connected to the new node.

4. The Consumer Process

4.1 Description

These processes wake up once every 30 seconds, and count the current number of nodes in the graph (which is in shared memory). Each consumer process gets mapped to a designated set of nodes which contains 1/10 of all nodes in the graph. Then each consumer process runs the Dijkstra's shortest path algorithms considering all nodes in its mapped set as the source nodes.

A consumer process runs Dijkstra's algorithm as stated above and writes (appends) the set of edges depicting the shortest path to each reachable node from the source node (along with the source node) to a file; then it goes to sleep, to wake up 30 seconds later and repeat the process.

4.2 Implementation

Each consumer is assigned 1/10 of the total number of nodes sequentially as its mapped-set. These consumers then each run the `dijkstra` function on their mapped-set and generate shortest path for each node in the graph. These paths are written to a file and the consumer process then goes to sleep for 30 seconds.

5. Optimization on Dijkstra

5.1 Description

Note that the producer will add only a limited number of nodes and/or edges to the (much larger) network, which might not change many shortest paths. Design and implement a strategy to optimize the process of re-computing shortest paths. This version will run with a “-optimize” flag to your code.

5.2 Implementation

Function prototypes

- void **dijkstra**(vector<int> sources, vector<int> &dist, vector<int> &parent) :
basic Dijkstra’s multi-source shortest path algorithm implemented on our Graph class
 - sources : vector of source indices for multi-source Dijkstra
 - dist : vector of distances of all nodes in the graph (initialised to INT_MAX)
 - parent : vector of immediate parent node of all nodes (initialised to -1)
 - void **optimized_dijkstra**(vector<int> sources, vector<int> &dist, vector<int> &parent)
(**&parent**)
 - sources : vector of source indices for multi-source Dijkstra
 - dist : vector of distances of all nodes in the graph, maintained for each consumer process
 - parent : vector of immediate parent node of all nodes, maintained for each consumer process
- Optimization :**
1. level-wise implementation of Dijkstra’s multi-source shortest path algorithm
 2. queue initialised with newly added source nodes
 3. all items in the queue processed in one go before moving to the next set of items in the queue (i.e. next level)
 4. early termination of function is possible if there are no updates to the dist vector for a particular level
- void **propagate_new_nodes**(vector<int> &dist, vector<int> &parent, vector<int> new_nodes)
(**new_nodes**)
 - dist : vector of distances obtained after calling optimized_dijkstra
 - parent : vector of immediate parent node of all nodes
 - new_nodes : non-source nodes newly added by the producer process

Propagation : We do the following process, starting with the new non-source nodes-

1. Nearest Neighbour Calculation - From the list of all neighbouring nodes, we find the node with least distance from any source node, we define this node as min_neighbouring_node and its corresponding distance as min_neighbouring_distance.
2. Distance and Parent Updation - We then update the distance of the present node as min_neighbouring_distance + 1 and set the parent as the min_neighbouring_node.

We then propagate to all neighbouring nodes and repeat this process.

Procedure : For each consumer process, $1/10th$ of the new nodes created by the producer are new source nodes whereas the rest are the new non-source nodes. In order to optimize the calculation of shortest paths, we do the following for each consumer process:

- At every iteration, we maintain the parent and the distance vector for all nodes which will be further used in the next iteration as the initial parent and distance vectors.
- We run the Dijkstra's multi-source algorithm in each consumer process with the source nodes as these new source nodes only. It will update the shortest path of a node only if it is shorter than the previously calculated path. This will eventually give us the shortest path of each older node to a source node.
- Further, in order to calculate the minimum distance of new non-source nodes, we run the `propagate_new_nodes` function which will calculate the shortest path of each new non-source node as described above.

These optimizations allow us to handle the new graph more efficiently, avoiding unnecessary computations. If the '`optimize`' flag is used, the faster `optimized_dijkstra` function is used on the new source set and the new non-source nodes have their distances updated using the `propagate_new_nodes` function.